

1. (50 points) The problem will focus on developing your own code for: *K-fold cross validation*.

Your code will be evaluated using five standard classification models applied to a multi-class classification dataset.

Dataset: We will be using the following dataset for the assignment.

Digits: The `Digits` dataset comes prepackaged with `scikit-learn` (`sklearn.datasets.load_digits`). The dataset has 1797 points, 64 features, and 10 classes corresponding to ten numbers $0, 1, \dots, 9$. The dataset was (likely) created from the following dataset:

<https://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits>

```
from sklearn.datasets import load_digits
digits = load_digits()
X, y = digits.data, digits.target
```

Classification Methods. We will consider five classification methods from `scikit-learn` and `xgboost`:

- Decision Tree Classifier: `DecisionTreeClassifier`,
- Naive Bayes Classifier: `GaussianNB`,
- Logistic Regression: `LogisticRegression`,
- Random Forest Classifier: `RandomForestClassifier`, and
- Gradient Boosting Classifier: `XGBClassifier`.

Use the following parameters for these methods (do not specify any additional parameters):

- `DecisionTreeClassifier`: `max_depth=10, random_state=42`
- `GaussianNB`: (no parameters needed)
- `LogisticRegression`: `penalty='l2', solver='lbfgs', random_state=42, multi_class='multinomial'`
- `RandomForestClassifier`: `max_depth=15, n_estimators=250, random_state=42`
- `XGBClassifier`: `max_depth=7, random_state=42`

You will have to develop `code` for the following three functions:

- (a) `get_model(method)`, which returns the appropriate classifier object with the parameters specified above, given the method specified as a string. The parameter `method` will be one of: `"DecisionTreeClassifier"`, `"GaussianNB"`, `"LogisticRegression"`, `"RandomForestClassifier"`, or `"XGBClassifier"`.
- (b) `get_splits(n, k, seed)`, which returns: randomized k ‘almost equal’ sized lists of unique disjoint indices from the set of all indices $\{0, \dots, n-1\}$, where the randomization depends on the integer `seed`. By ‘almost equal’, we mean the cardinality of the lists can differ by at most 1. These k list of indices correspond to the k folds over which cross-validation will be performed. The function will have the following `output`:
 - i. a python list containing exactly k lists. Each of these sublists should be disjoint, each of size roughly $\frac{n}{k}$, contain elements from the set $\{0, 1, \dots, n-1\}$ and must not contain repeated elements. The union of all the k sublists should include all elements in $\{0, \dots, n-1\}$.

Input `seed` determines the randomization and the output should be the same every time we use the same `seed` for a given n, k , and should be (randomly) different for different values of the `seed`.

For example, `get_splits(4, 2, 1)` may return `[[0,2], [1,3]]`, and the output must be same every time with the same input; `get_splits(4, 2, 73)` may return `[[0,3], [1,2]]`; `get_splits(7, 2, 2)` may return `[[0,2,4,6], [1,3,5]]`; `get_splits(11, 3, 7)` may return `[[0,3,6,9], [1,4,7,10], [2,5,8]]`.

For our tests, n would be less than 1200. `get_splits` would have a time limit of 10 seconds.

- (c) `my_cross_val(method, X, y, splits)`, which runs k -fold cross-validation for `method` on the dataset (X, y) . The **input parameters** are:

- i. `method`, which specifies the (class) name of one of the five classification methods under consideration,
- ii. `X, y` which is the data for the classification problem
- iii. `splits`: the output of the `get_splits` method (`len(splits) = k`)

The function will have the following **output**:

- i. the test set error rates for each of the k folds.

The error should be measured as $\frac{\# \text{ of wrong predictions}}{\# \text{ of total predictions}}$. The (auto)grader will judge your solution as correct if the difference between the reported and the expected mean error rates is within 10^{-3} . `my_cross_val` would have a time limit of 2 minutes.

Within `my_cross_val`, strictly use the `splits` encoded in the input parameter `splits`. Do not define your `splits` for K -fold cross validation within this method.

Also, make sure that you are NOT inadvertently shuffling your training data during K -fold cross validation. The training examples (for any particular split) should be in the same order as the input X .

Use `my_cross_val` to return the error rates in each fold for k fold cross-validation for the specific `method` (one of `LinearSVC`, `SVC`, `LogisticRegression`, `RandomForestClassifier`, and `XGBClassifier`) with the parameters outlined above.

2. (50 points) The assignment will focus on developing your own code for: `random train-test split validation`.

Your code will be evaluated using five standard classification models applied to a multi-class classification dataset.

Dataset: We will be using the following dataset for the assignment.

Digits: The Digits dataset comes prepackaged with `scikit-learn` (`sklearn.datasets.load_digits`). The dataset has 1797 points, 64 features, and 10 classes corresponding to ten numbers $0, 1, \dots, 9$. The dataset was (likely) created from the following dataset:

<https://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits>

Classification Methods. We will consider five classification methods from `scikit-learn` and `xgboost`:

- Decision Tree Classifier: `DecisionTreeClassifier`,
- Naive Bayes Classifier: `GaussianNB`,
- Logistic Regression: `LogisticRegression`,
- Random Forest Classifier: `RandomForestClassifier`, and
- Gradient Boosting Classifier: `XGBClassifier`.

Use the following parameters for these methods (do not specify any additional parameters):

- `DecisionTreeClassifier`: `max_depth=10, random_state=42`
- `GaussianNB`: (no parameters needed)
- `LogisticRegression`: `penalty='l2', solver='lbfgs', random_state=42, multi_class='multinomial'`
- `RandomForestClassifier`: `max_depth=15, n_estimators=250, random_state=42`
- `XGBClassifier`: `max_depth=7, random_state=42`

Your `homework2_q2.py` file should include and make use of the same `get_model(method)` function as in `homework2_q1.py`

Develop code for `my_train_test(method, X, y, π, k)`, which performs random splits on the data (X, y) so that $π \in [0, 1]$ fraction of the data is used for training using `method`, rest is used for testing, and the process is repeated k times, after which the code returns the error rate for each such train-test split. Your `my_train_test` will be tested with $π = 0.75$ and $k = 10$ on the five methods above applied to the `Digits` dataset.

You will have to develop new code for the following function:

`my_train_test(method, X, y, π, k)`, which does random train-test split based evaluation of `method` with $π$ fraction used for training for each split.

The function will have the following **input**:

- (1) `method`, which specifies the (class) name of one of the five classification methods under consideration,
- (2) X, y which is data for the classification problem,
- (3) $π$, the fraction of data chosen randomly to be used for training,

(4) k , the number of times the train-test split will be repeated.

The function will have the following **output**:

- (a) A list of the test set error rates for each of the k splits.

Error rate should be calculated as $\frac{\# \text{ of wrong predictions}}{\# \text{ of total predictions}}$. The (auto)grader will compare the mean and standard deviation of your list with our solution; it must be within three standard deviations. `my_train_test` would have a time limit of 2 minutes.